

# PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation

李声涛 17214643 宋日辉 17214675  
黄莉 17214619 管卓群 17214612  
周晓梅 17214729

School of Data and Computer Science, SYSU

December 19, 2017

# Content


- Introduction
- Problem Statement
- Deep Learning on Point Sets
- Experiment
- Conclusion

# Introduction

- *Title:*
  - PointNet: Deep Learning on Point Sets for 3D Classification and Segmentation
- *Keywords:*
  - Deep Learning
  - Point Sets (3D)
  - Classification and Segmentation
  - PointNet

# Deep Learning


- Part of a broader family of machine learning methods
- Based on learning data representations
- Deep learning architectures: DNN, DBN, RNN, CNN...




(a) Deep Learning Structure

# Point Sets

- A set of points.
- 2D point:  $(x, y)$
- 3D point:  $(x, y, z)$




(b) 2D Point Set



(c) 3D Point Set (Point Cloud)

# Classification and Segmentation


- Classification: distinguish different classes
- Segmentation: figure out different parts



(d) Room

# Neural Network Based Methods


- Volumetric CNNs: 3D voxel grids
  - Constrained by resolution



(e) 3D ShapeNet

# Neural Network Based Methods

- Multi-view CNNs: collections of images
  - Nontrivial to extend them to scene understanding or other 3D tasks.



(f) MVCNN



# PointNet

- A novel deep net architecture
- Input: points
- Output: points and scores for each label
- Tasks: 3D shape classification, shape part segmentation, and scene semantic parsing
- Simple, effective and robust

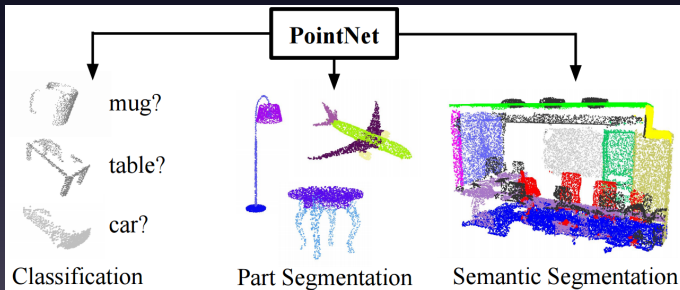


Figure: Task of PointNet

# Problem Statement

- A point cloud is represented as a set of 3D points:
  - $\{P_i | i = 1, \dots, n\}$
  - $P_i = (x, y, z)$
  - Extra feature channels: color, normal, etc.

# Object Classification

- The input:  
Directly sampled from a shape  
Pre-segmented from a scene point cloud.
- The output:  
This deep network outputs  $k$  scores for all the  $k$  candidate classes.

# Semantic Segmentation

- The input:  
A single object for part region segmentation  
A sub-volume from a 3D scene for object region segmentation.
- The output:  
This model will output  $n \times m$  scores for each of the  $n$  points and each of the  $m$  semantic subcategories.

# Deep Learning on Point Sets

- Properties of Point Sets
- PointNet Architecture

# Properties of Point Sets

- Unordered
- Interaction among points
- Invariance under transformations

# PointNet Architecture

- Symmetry Function for Unordered Input
- Local and Global Information Aggregation
- Joint Alignment Network




Figure: Architecture

# Experiments

- Applications
- Architecture Design Analysis
- Visualizing PointNet
- Time and Space Complexity Analysis



# Applications-3D Object Classification

- 12,311 CAD models
- from 40 man-made object categories,
- split into 9,843 for training and 2,468 for testing

	input	#views	accuracy avg. class	accuracy overall
SPH [11]	mesh	-	68.2	-
3DShapeNets [25]	volume	1	77.3	84.7
VoxNet [15]	volume	12	83.0	85.9
Subvolume [16]	volume	20	86.0	<b>89.2</b>
LFD [25]	image	10	75.5	-
MVCNN [20]	image	80	<b>90.1</b>	-
Ours baseline	point	-	72.6	77.4
Ours PointNet	point	1	86.2	<b>89.2</b>

Table 1. **Classification results on ModelNet40.** Our net achieves state-of-the-art among deep nets on 3D input.

# Applications-3D Object Part Segmentation

- ShapeNet part data set
- 16,881 shapes from 16 categories, annotated with 50 parts in total
- mIoU?

	mean	aero	bag	cap	car	chair	ear phone	guitar	knife	lamp	laptop	motor	mug	pistol	rocket	skate board	table
# shapes		2690	76	55	898	3758	69	787	392	1547	451	202	184	283	66	152	5271
Wu [24]	-	63.2	-	-	-	73.5	-	-	-	74.4	-	-	-	-	-	-	74.8
Yi [26]	81.4	81.0	78.4	77.7	<b>75.7</b>	87.6	61.9	<b>92.0</b>	85.4	<b>82.5</b>	<b>95.7</b>	<b>70.6</b>	91.9	<b>85.9</b>	53.1	69.8	75.3
3DCNN	79.4	75.1	72.8	73.3	70.0	87.2	63.5	88.4	79.6	74.4	93.9	58.7	91.8	76.4	51.2	65.3	77.1
Ours	<b>83.7</b>	<b>83.4</b>	<b>78.7</b>	<b>82.5</b>	74.9	<b>89.6</b>	<b>73.0</b>	91.5	<b>85.9</b>	80.8	95.3	65.2	<b>93.0</b>	81.2	<b>57.9</b>	<b>72.8</b>	<b>80.6</b>

Table 2. **Segmentation results on ShapeNet part dataset.** Metric is mIoU(%) on points. We compare with two traditional methods [24] and [26] and a 3D fully convolutional network baseline proposed by us. Our PointNet method achieved the state-of-the-art in mIoU.

# Applications-Semantic Segmentation in Scenes

- Stanford 3D semantic parsing data set
- The dataset contains 3D scans from Matterport scanners in 6 areas including 271 rooms. Each point in the scan is annotated with one of the semantic labels from 13 categories (chair, table, floor, wall etc)



# Applications-Semantic Segmentation in Scenes

	mean IoU	overall accuracy
Ours baseline	20.12	53.19
Ours PointNet	<b>47.71</b>	<b>78.62</b>

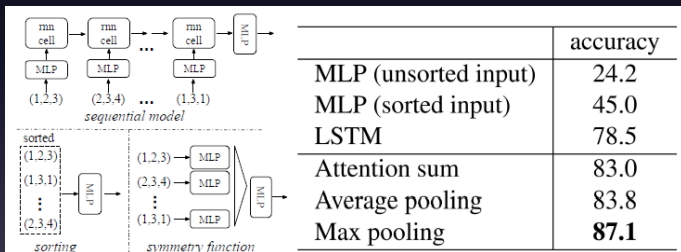
Table 3. **Results on semantic segmentation in scenes.** Metric is average IoU over 13 classes (structural and furniture elements plus clutter) and classification accuracy calculated on points.

	table	chair	sofa	board	mean
# instance	455	1363	55	137	
Armeni et al. [1]	46.02	16.15	<b>6.78</b>	3.91	18.22
Ours	<b>46.67</b>	<b>33.80</b>	4.76	<b>11.72</b>	<b>24.24</b>

Table 4. **Results on 3D object detection in scenes.** Metric is average precision with threshold IoU 0.5 computed in 3D volumes.

# Architecture Design Analysis

- Three approaches to achieve order invariance.
- ModelNet40 shape classification problem




# Architecture Design Analysis

- Effectiveness of Input and Feature Transformations
- ModelNet40 shape classification problem

Transform	accuracy
none	87.1
input (3x3)	87.9
feature (64x64)	86.9
feature (64x64) + reg.	87.4
both	<b>89.2</b>


# Architecture Design Analysis

- Robustness Test
- ModelNet40 shape classification problem



# Visualizing PointNet

- critical point sets  $C_S$  & the upper-bound shapes  $N_S$





# Time and Space Complexity Analysis

- PointNet's space and time, complexity is  $O(N)$
- point cloud classification: 1K objects/second
- semantic segmentation: 2 rooms/second
- 1080X GPU on TensorFlow


	#params	FLOPs/sample
PointNet (vanilla)	0.8M	148M
PointNet	3.5M	440M
Subvolume [16]	16.6M	3633M
MVCNN [20]	60.0M	62057M

# Conclusion

- A brief introduction
- PointNet architecture
- Experiment result

# Repeat the experiment

- Tensorflow
- CPU: i7 - 5700
- GPU: Geforce 1070
- Training time: 2h31min(classification) and 16h28min(part segmentation)




# Heuristics


- Use it on robots
- Learn to create a novel structure
- Figure out a research idea

# Traditional Point Cloud Processing

- Edge-based methods
- Model-based methods
- Region-based methods
- Attributes-based methods
- Graph-based methods



(a) A cup on the desk



(b) A part of leg